

Gaussian Process

Jiali Lin

Virginia Tech

December 4, 2016

Outline

Introduction

GPs for regression

GPs meet GLMs

Introduction

- ▶ We observe some inputs \mathbf{x}_i and some outputs y_i (i.i.d).
- ▶ We assume linear relationships

$$y = f(\mathbf{x}) = \mathbf{x}^T \mathbf{w}, \quad y = f(\mathbf{x}) + \epsilon, \quad \epsilon \sim N(0, \sigma_n^2)$$

- ▶ Prior: $\mathbf{w} \sim N(0, \mathbf{\Sigma}_p)$.
- ▶ To make predictions on new input \mathbf{x}_* , we use posterior predictive distribution

$$p(y_* | \mathbf{x}_*, \mathbf{X}, \mathbf{y}) = \int p(y_* | f, \mathbf{x}_*) p(f | \mathbf{X}, \mathbf{y}) df$$

- ▶ Alternatively one could first map \mathbf{x} to some basis function, then let

$$f(\mathbf{x}) = \phi(\mathbf{x})^T \mathbf{w}$$

Main Ideas

- ▶ Want more flexible form for $f(x)$, treat the whole $f(\cdot)$ as a parameter.
- ▶ Assume $f(\cdot)$ takes values from a function space.
- ▶ Assume $f(\cdot)$ is random. In particular, Gaussian Process.
- ▶ **Gaussian processes** or **GPs**: defines a prior over functions, which can be converted into a posterior over functions once we have seen some data.

Outline

Introduction

GPs for regression

GPs meet GLMs

Gaussian process

- ▶ **Definition:** Gaussian process is a collection of random variables, any *finite* number of which have a joint Gaussian distribution.
- ▶ Let the prior on the regression function be a GP, denoted by

$$f(\mathbf{x}) \sim \text{GP}(m(\mathbf{x}), \kappa(\mathbf{x}, \mathbf{x}'))$$

- ▶ $m(\mathbf{x})$ is the mean function and $\kappa(\mathbf{x}, \mathbf{x}')$ is the covariance function

$$m(\mathbf{x}) = E[f(\mathbf{x})]$$

$$\kappa(\mathbf{x}, \mathbf{x}') = E[(f(\mathbf{x}) - m(\mathbf{x}))(f(\mathbf{x}') - m(\mathbf{x}'))^T]$$

- ▶ Require $\kappa()$ be a positive definite kernel. For any finite set of points, this process defines a joint Gaussian

$$p(\mathbf{f}|\mathbf{X}) = N(\mathbf{f}|\boldsymbol{\mu}, \mathbf{K})$$

- ▶ $K_{ij} = \kappa(\mathbf{x}_i, \mathbf{x}_j)$ and $\boldsymbol{\mu} = (m(\mathbf{x}_1), \dots, m(\mathbf{x}_N))$. Note that it is common to use a mean function of $m(\mathbf{x}) = 0$, since the GP is flexible enough to model the mean arbitrarily well.

Predictions using noise-free observations

- Given noise-free training data

$$\mathcal{D} = \{\mathbf{x}^{(i)}, f^{(i)} | i = 1, \dots, n\}$$

- We want to make predictions \mathbf{f}_* at test points \mathbf{X}_* .
- By definition of the GP, the joint distribution has the following form

$$\begin{bmatrix} \mathbf{f} \\ \mathbf{f}_* \end{bmatrix} \sim N \left(\begin{bmatrix} \boldsymbol{\mu} \\ \boldsymbol{\mu}_* \end{bmatrix}, \begin{bmatrix} K(\mathbf{X}, \mathbf{X}) & K(\mathbf{X}, \mathbf{X}_*) \\ K(\mathbf{X}_*, \mathbf{X}) & K(\mathbf{X}_*, \mathbf{X}_*) \end{bmatrix} \right)$$

- Posterior: $p(\mathbf{f}_* | \mathbf{X}_*, \mathbf{X}, \mathbf{f}) \sim N(\mathbf{f}_* | \boldsymbol{\mu}_*, \boldsymbol{\Sigma}_*)$

$$\begin{aligned} \boldsymbol{\mu}_* &= \boldsymbol{\mu}(\mathbf{X}_*) + K(\mathbf{X}_*, \mathbf{X})K(\mathbf{X}, \mathbf{X})^{-1}(\mathbf{f} - \boldsymbol{\mu}(\mathbf{X})) \\ \boldsymbol{\Sigma}_* &= K(\mathbf{X}_*, \mathbf{X}_*) - K(\mathbf{X}, \mathbf{X}_*)K(\mathbf{X}, \mathbf{X})^{-1}K(\mathbf{X}_*, \mathbf{X}) \end{aligned}$$

Predictions using noise-free observations (cont'd)




Figure: Left: some functions sampled from a GP prior with SE kernel. Right: some samples from a GP posterior, after conditioning on 5 noise-free observations. The shaded area represents $E[f(x)] \pm 2\text{std}(f(x))$. Based on Figure 2.2 of (Rasmussen and Williams 2006). Figure generated by GprDemo.

Predictions using noisy observations

- ▶ $y = f(\mathbf{x}) + \epsilon$, where $\epsilon \sim N(0, \Sigma_y^2)$.
- ▶ The covariance of the observed noisy responses is

$$\text{cov}[y_p, y_q] = \kappa(\mathbf{x}_p, \mathbf{x}_q) + \Sigma_y^2 \delta_{pq}, \quad \text{where} \quad \delta_{pq} = \mathbb{I}(p = q)$$

- ▶ We assume the noise terms were independent.
- ▶ Set mean function $m(\mathbf{x}) = 0$. Thus, $f(\mathbf{x}) \sim \text{GP}(0, \kappa(\mathbf{x}, \mathbf{x}'))$
- ▶ Now the joint distribution is given by

$$\begin{bmatrix} \mathbf{y} \\ \mathbf{f}_* \end{bmatrix} \sim N \left(\begin{bmatrix} \boldsymbol{\mu} \\ \boldsymbol{\mu}_* \end{bmatrix}, \begin{bmatrix} K(\mathbf{X}, \mathbf{X}) + \sigma_n^2 \mathbf{I} & K(\mathbf{X}, \mathbf{X}_*) \\ K(\mathbf{X}_*, \mathbf{X}) & K(\mathbf{X}_*, \mathbf{X}_*) \end{bmatrix} \right)$$

- ▶ Posterior: $p(\mathbf{f}_* | \mathbf{X}_*, \mathbf{X}, \mathbf{f}) \sim N(\mathbf{f}_* | \boldsymbol{\mu}_*, \boldsymbol{\Sigma}_*)$

$$\boldsymbol{\mu}_* = K(\mathbf{X}_*, \mathbf{X})[K(\mathbf{X}, \mathbf{X}) + \sigma_n^2 \mathbf{I}]^{-1} \mathbf{y}$$

$$\boldsymbol{\Sigma}_* = K(\mathbf{X}_*, \mathbf{X}_*) - K(\mathbf{X}, \mathbf{X}_*)[K(\mathbf{X}, \mathbf{X}) + \sigma_n^2 \mathbf{I}]^{-1} K(\mathbf{X}_*, \mathbf{X})$$

Effect of the kernel parameters

- ▶ The predictive performance of GPs depends exclusively on the chosen kernel.
- ▶ Suppose we choose **squared-exponential** (SE) kernel for the noisy observations

$$\kappa_y(x_p, x_q) = \sigma_f^2 \exp\left(-\frac{1}{2\ell^2}(x_p - x_q)^2\right) + \sigma_y^2 \delta_{pq}$$

- ▶ ℓ is the horizontal scale over which the function changes, σ_f^2 controls the vertical scale of the function, and σ_y^2 is the noise variance.

Effect of the kernel parameters (cont'd)




Figure: Some 1d GPs with SE kernels but different hyper-parameters fit to 20 noisy observations. The kernel has the form in Equation (14). The hyper-parameters $(\ell, \sigma_f, \sigma_y)$ are as follows: (a) $(1, 1, 0.1)$ (b) $(0.3, 0.108, 0.00005)$, (c) $(3.0, 1.16, 0.89)$. Figure generated by `gprDemoChangeHparams`, written by Carl Rasmussen.

- ▶ In Figure (a), the result is a good fit.
- ▶ In Figure (b), the function looks more “wiggly”. Also, the uncertainty goes up faster.
- ▶ In Figure (c), the function looks smoother.

Estimating the kernel parameters

- ▶ Frequentist: exhaustive search over a discrete grid of values (slow!).
- ▶ Consider empirical Bayes approach.
- ▶ Maximization log likelihood can be done using efficient gradient-based optimization algorithms.
- ▶ In absence of prior $p(\boldsymbol{\theta})$, the posterior for hyperparameter $\boldsymbol{\theta}$ is proportional to the marginal likelihood

$$p(\boldsymbol{\theta}|\mathbf{X}, \mathbf{y}) \propto p(\mathbf{y}|\mathbf{X}, \boldsymbol{\theta})$$

- ▶ Choose $\boldsymbol{\theta}$ to optimize the marginal log-likelihood

$$\log p(\boldsymbol{\theta}|\mathbf{X}, \mathbf{y}) \propto -\frac{1}{2} \log |K(\mathbf{X}, \mathbf{X}) + \sigma^2 \mathbf{I}| - \frac{1}{2} \mathbf{y}^T (K(\mathbf{X}, \mathbf{X}) + \sigma^2 \mathbf{I})^{-1} \mathbf{y}$$

Outline

Introduction

GPs for regression

GPs meet GLMs

Gaussian Process Classification

- In the binary case, we have

$$y_i \in \{0, 1\}$$

$$p(y_i | \mathbf{x}_i, f_i) = \exp\{y_i f_i - A(f_i)\}$$

$$p(\mathbf{y}) = N(\mathbf{f} | \mathbf{0}, \mathbf{K}), \quad \text{where} \quad K_{ij} = \kappa(\mathbf{x}_i, \mathbf{x}_j)$$

$$p(y_i | \mathbf{x}_i, f_i) = \text{Ber}(y_i | \text{Sigm}(f_i))$$

- Equivalent to logistic regression, but uses kernels rather than features.
- Gaussian prior on weights replaced by Gaussian prior on training log-odds.
- Basic inference finds MAP estimate of function f at all training points

$$\hat{\mathbf{f}} = \underset{\mathbf{f}}{\operatorname{argmax}} \log p(\mathbf{f}) + \sum_{i=1}^N \log p(y_i | \mathbf{x}_i, f_i)$$

Gaussian Process Classification (cont'd)

- Interpretation of function values f_i
 - **Positive:** $p(y_i = 1|f_i) > 0.5$.
 - **Zero:** $p(y_i = 1|f_i) = 0.5$.
 - **Negative:** $p(y_i = 1|f_i) < 0.5$.
- Interpretation of kernel values K_{ij}
 - **Positive:** likely have same label.
 - **Zero:** inputs are totally independent.
 - **Negative:** likely have different labels.

Gaussian Process Classification (cont'd)

First, compute the distribution of the latent variable for a test case

$$p(f_*|\mathbf{x}_*, \mathbf{X}, \mathbf{y}) = N(\mathbb{E}[f_*], \text{var}[f_*])$$

Second, produce a predictive distribution for binary responses

$$\pi_* = p(y_* = 1|\mathbf{x}_*, \mathbf{X}, \mathbf{y}) \approx \int \sigma(f_*)p(f_*|\mathbf{x}_*, \mathbf{X}, \mathbf{y})df_*$$

- ▶ Note that $p(\mathbf{y}|\mathbf{f})$ has link function involved, conjugacy of \mathbf{f} are lost.
- ▶ Also integrations are difficult.
- ▶ **Solutions:** use analytic approximations of integrals to approximate $p(\mathbf{f}|\mathbf{X}, \mathbf{y})$, or Monte Carlo sampling.